


Foundations of High Speed Cryptography

Module 1 - Theory

Lesson 3 - Cryptographic Hashes and Merkle Trees

Karthik Inbasekar



Agenda

- **Motivation**
- **Cryptographic Hash Functions (CHF)**
 - Collision Resistant Hash Functions: Definitions and examples
 - Bitwise hashes: Sponge Construction, eg: SHA3
 - Finite Field friendly Hashes: General structure, eg: Poseidon2
- **Merkle Tree**
 - General principles
 - Merkle commitment Scheme
 - Optimizations
- **Hash Functions as Random Oracles**
 - Fiat-Shamir transformation

Motivation

- Cryptographic hash: **fundamental cryptographic primitive** that enables mapping any input into a fixed-size digest that “looks random”.

Cryptographic hash functions have many information-security applications, notably in digital signatures, message authentication codes (MACs), and other forms of authentication. They can also be used as ordinary hash functions, to index data in hash tables, for fingerprinting, to detect duplicate data or uniquely identify files, and as checksums to detect accidental data corruption. Indeed, in information-security contexts, cryptographic hash values are sometimes called (digital) fingerprints, checksums, (message) digests, or just hash values, even though all these terms stand for more general functions with rather different properties and purposes. Non-cryptographic hash functions are used in hash tables and to detect accidental errors; their constructions frequently provide no resistance to a deliberate attack. For example, a denial-of-service attack on hash tables is possible if the collisions are easy to find, as in the case of linear cyclic redundancy check (CRC) functions

1043 bytes

Hash

32 bytes

6ccbbefe857c8fd295e580eb71
2f0e7de4aba31a663b50374cf92
1024335e0e3

- **Useful building blocks** : Data integrity, Consensus: proof of work , Cryptographic commitment Schemes, Pseudo Random Number Generators, digital signatures, password verification
- Digests do not reveal actual contents of data, can be used for **validation of information while preserving privacy**.

Use Cases

- **Data integrity:** unique fingerprints of blocks/transactions in Ethereum

Keccak : Account, Storage, Transaction and receipt Tries.

- **Proof of work:** Consensus mechanism in Bitcoin

Double SHA-256 for hashing, allows decentralized agreement on transaction history.

- **Cryptographic Commitment Schemes:** Hashes can act as binding and hiding commitments

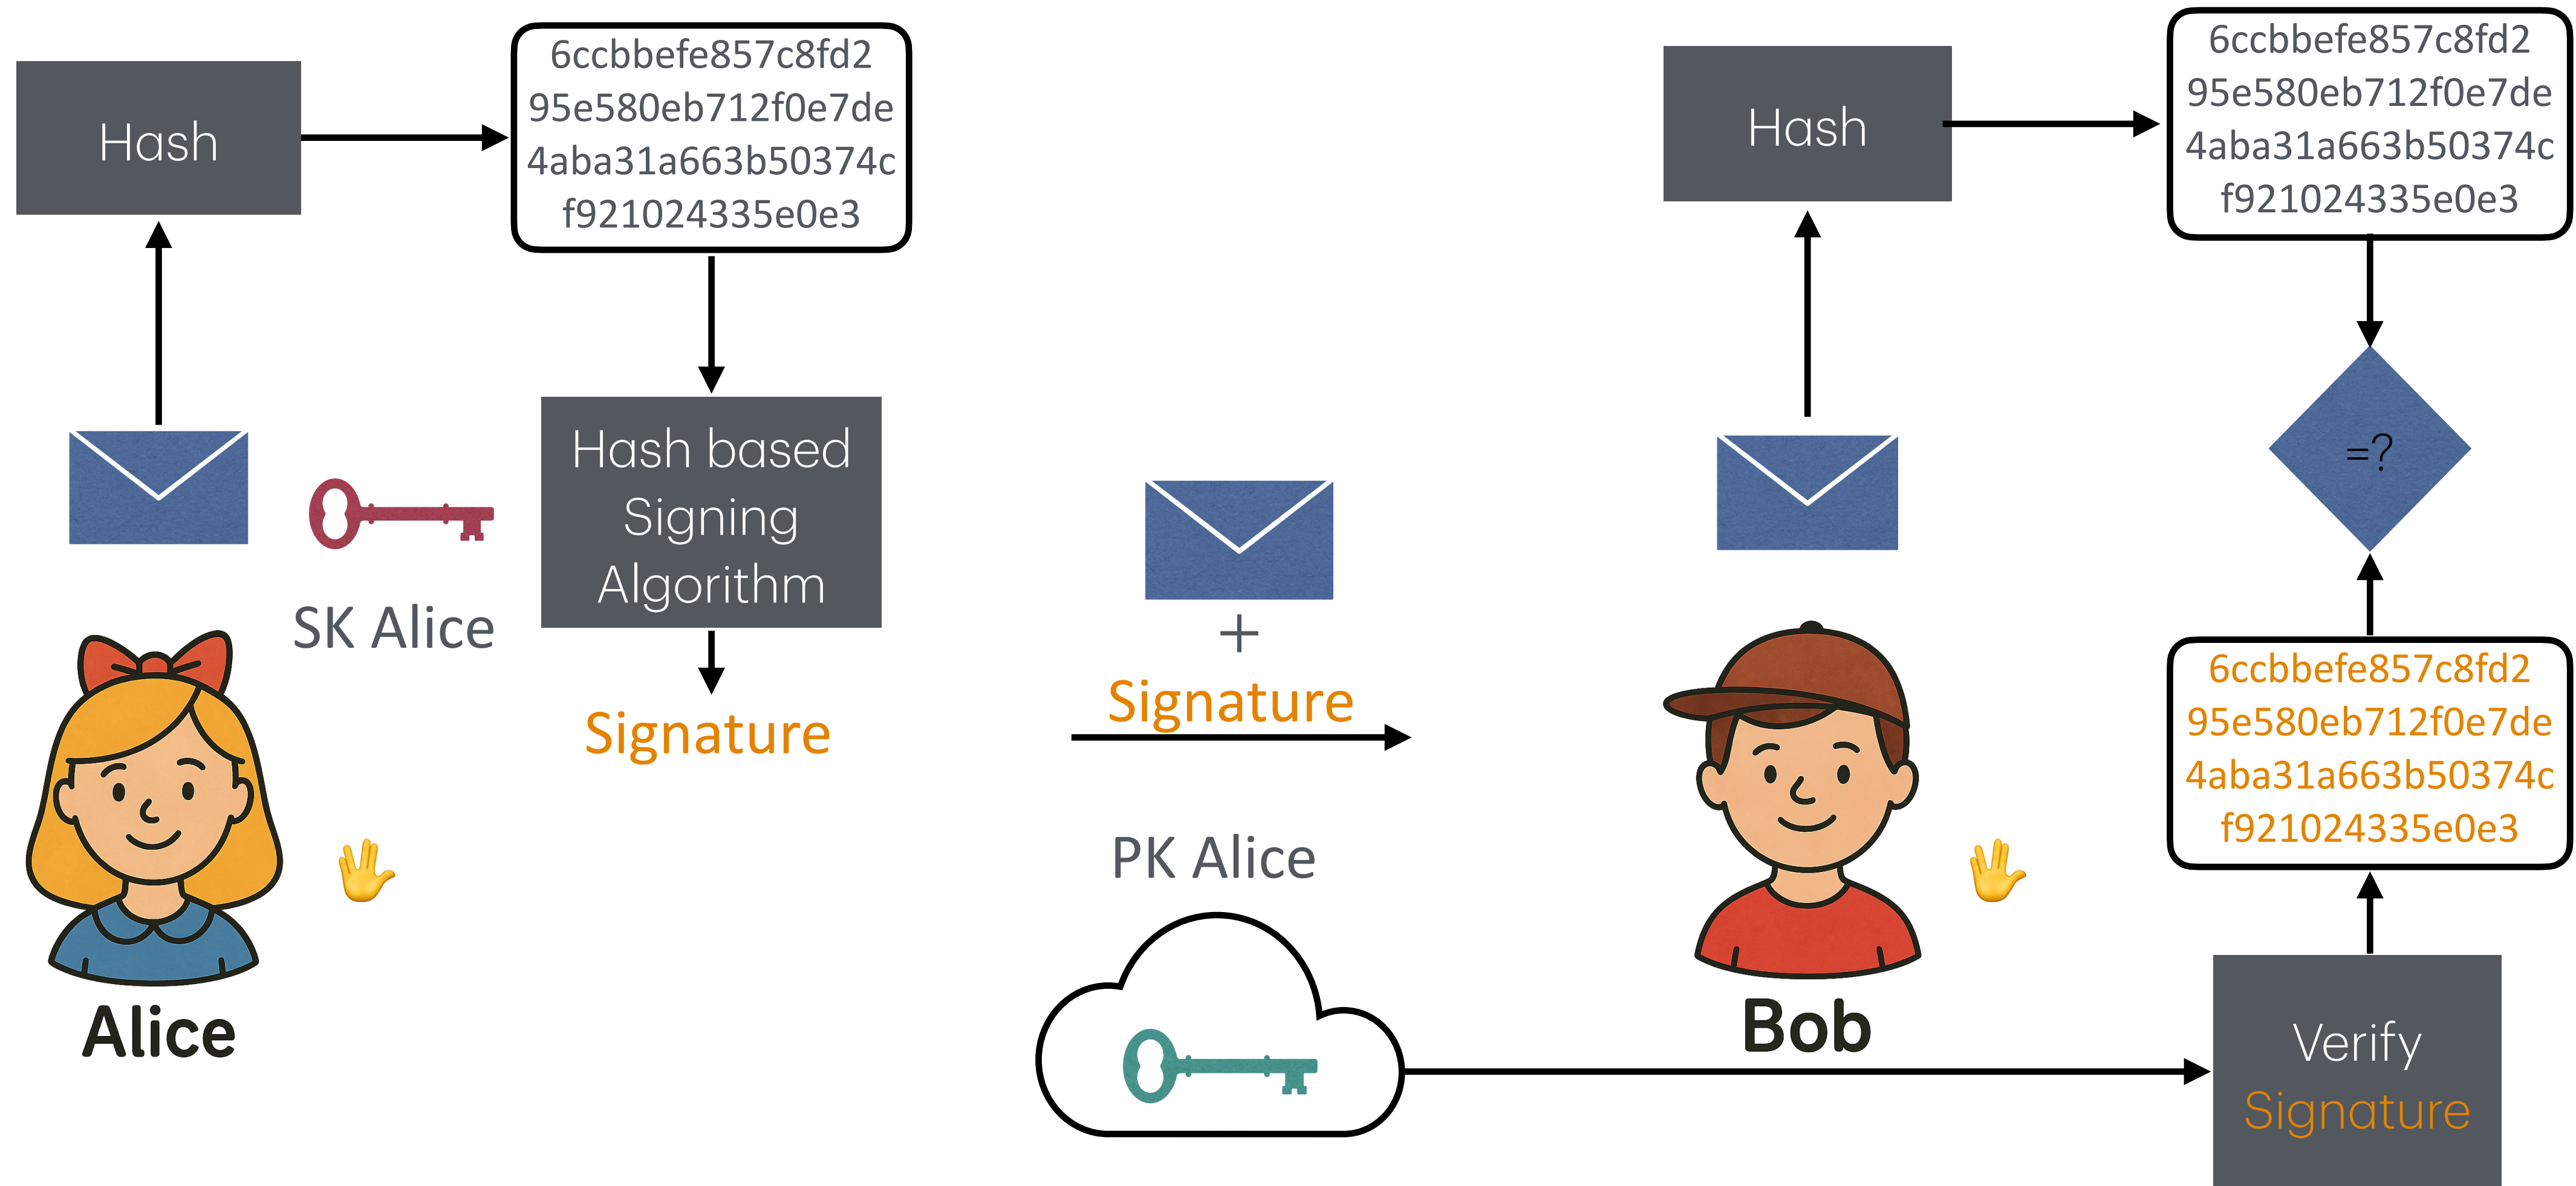
Merkle proofs : proving knowledge of data, used in Cryptographic proof systems.

- **Fiat Shamir Heuristic:** Turn interactive proof systems into non-interactive proof system.

- **Digital Signatures:** Post quantum digital signatures, are built on Cryptographic Hash Functions.

Eg: XMSS, SPHINCS++

Digital Signatures using hash functions




- Eg: XMSS , SPHINCS++

<https://csrc.nist.gov/csrc/media/events/workshop-on-cybersecurity-in-a-post-quantum-world/documents/papers/session5-hulsing-paper.pdf>

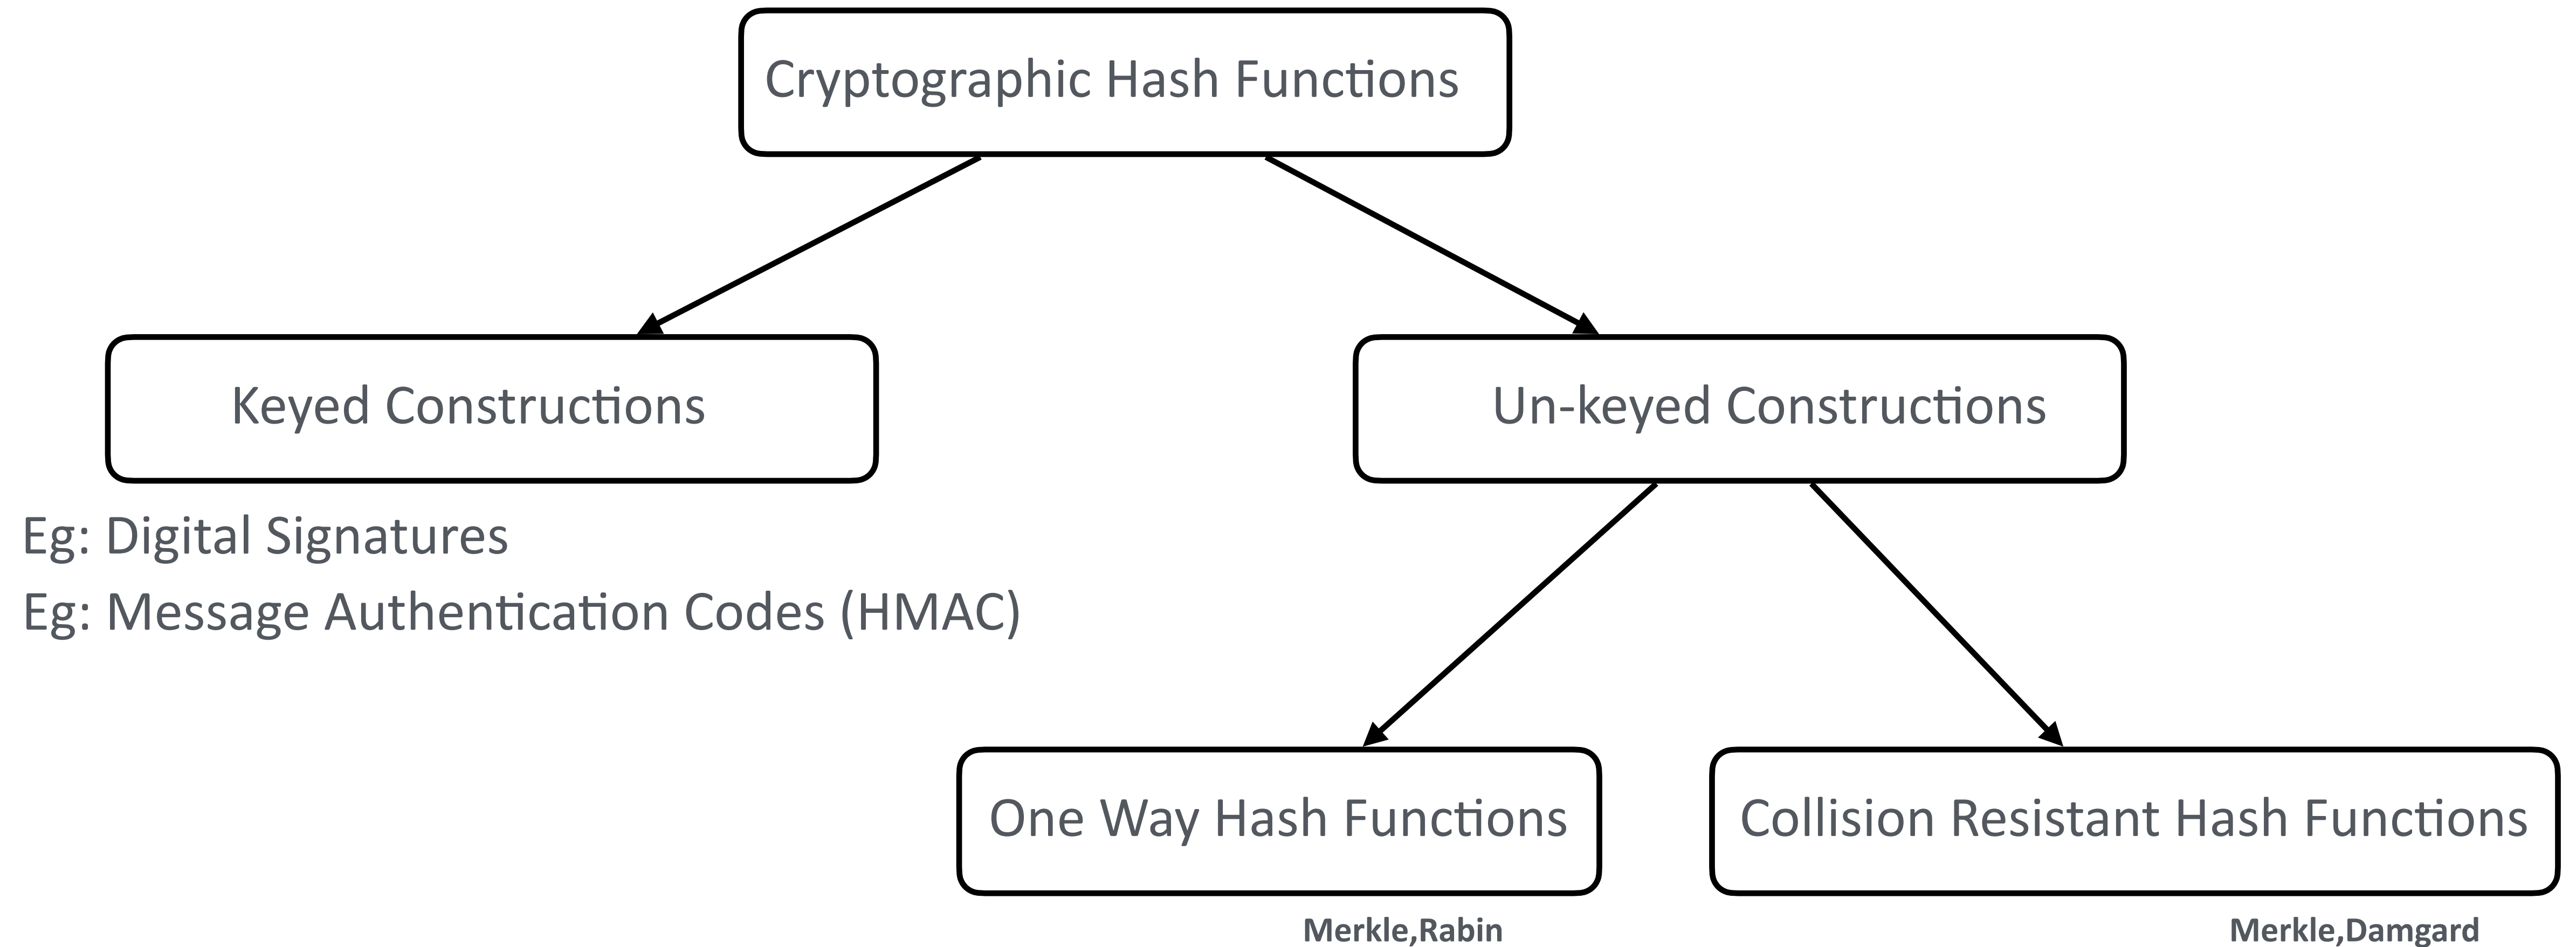
Foundations of High Speed Cryptography

Module 1 - Theory
Lesson 3 - Hashes and Merkle Trees



Cryptographic Hash functions

Taxonomy



Collision Resistant Hash Functions

Given a hash h , the argument x can be of arbitrary length and the digest $y = h(x)$ of length $n \geq 256$

- **Pre-image resistance**

Given y, h computationally infeasible to find $x \mid h(x) = y$

- **Second Pre-image resistance**

Given $x, y = h(x)$ it is computationally infeasible to find $x' \mid h(x') = h(x)$

- **Collision resistance**

It is computationally feasible to find any two distinct inputs $x \neq x' \mid h(x') = h(x)$

Birthday attack $\sim 2^{n/2}$

- **Pseudo randomness**

The output of the hash function must not have any detectable output patterns that leak information.

Preimage resistance vs Collision resistance

- Given $SHA256(x)$ is a known collision resistant function: what about this one?

$$h(x) = SHA256(x) \bmod 2^{32}$$

- Given $g(x)$ is collision resistant, and with digest size $n - 1$ bits

$$h(x) = \begin{cases} 1 || x, & \text{if } \text{len}(x) = n - 1 \\ 0 || g(x), & \text{otherwise} \end{cases}$$


- Given $h(x) = x \forall x \in \{0,1\}^n$

- Given $g(x)$ is collision resistant (by some hard problem LWE/SIS/DLOG) such that

$$g(x_1) + g(x_2) = g(x_1 + x_2) \quad \text{Eg SWIFFT}$$

Bitwise hashes: SHA -3



- SHA3 uses the sponge construction to hash arbitrary length input to a fixed size output
- The hash state has a public part r : **rate** and a private part c : **capacity**, $r + c = 1600$ bits
- The input is chunked into $m_0 || m_1 || m_2 || \dots || m_k ||$ such that $|m_i| = r$ bits



- The permute function P consists of multiple rounds, input chunks are XORed to rate, output chunks are read r bits at a time.

SHA-3 : Keccak-f permutation

- The hash state is represented as an array $5 \times 5 \times 64 = 1600$




25 lanes each of 64 bit word size

SHA-3 : Keccak-f permutation

- Each permute operation operates for $12 + 2l$ rounds in the five step sequence

Step	Operation	Optimization
\oplus	For each bit, XOR with parity of neighboring columns	Parallelize : Column parity, XOR of bits with parity
P	Rotate each bit of a lane by an offset length	Parallelize: Bit rotation per lane
Π	Permute the positions of the Lanes	Pure reordering, parallel across all lanes
X	XOR each bit with a non linear function of two other bits in a row	Row wise parallel
I	Add round constants to the lane $A[0,0,Z]$	Store round constants in a cached read only memory



- Bit interleaving for area/speed tradeoffs: $25s$ words of $64/s$ bits for different hardware platforms.

Finite Field Hashes

- In cryptographic constructions involving Proofs of computational integrity, one is required to prove that a given hash state is valid, and the hash was executed correctly
- In such constructions, each bitwise operation of the hash is represented as an arithmetic circuit.




$$Q(A, B) = A + B - 2 \cdot A \cdot B$$

XOR Constraint equation

- For 64kb Sha256 we need to check about 30 M - 45 M constraint equations!!
- Due to large number of bitwise operations, the arithmetic circuit representation consists of large number of gates, and provable constraints that increase prover overhead significantly,
- Finite Field hashes, operate on finite fields instead of bits, and have a much simpler arithmetic circuit representation.

Finite Field Hashes

- A popular construction is based on the HADES family of block ciphers



- The state is a vector of size t , with elements in \mathbb{F}_p^t
- Cipher operations

- Linear layer: Add round constant vectors
- Linear layer: Matrix multiplications by an MDS matrix

- Full Non-Linear layers:

$$\gcd(\alpha, p - 1) = 1$$

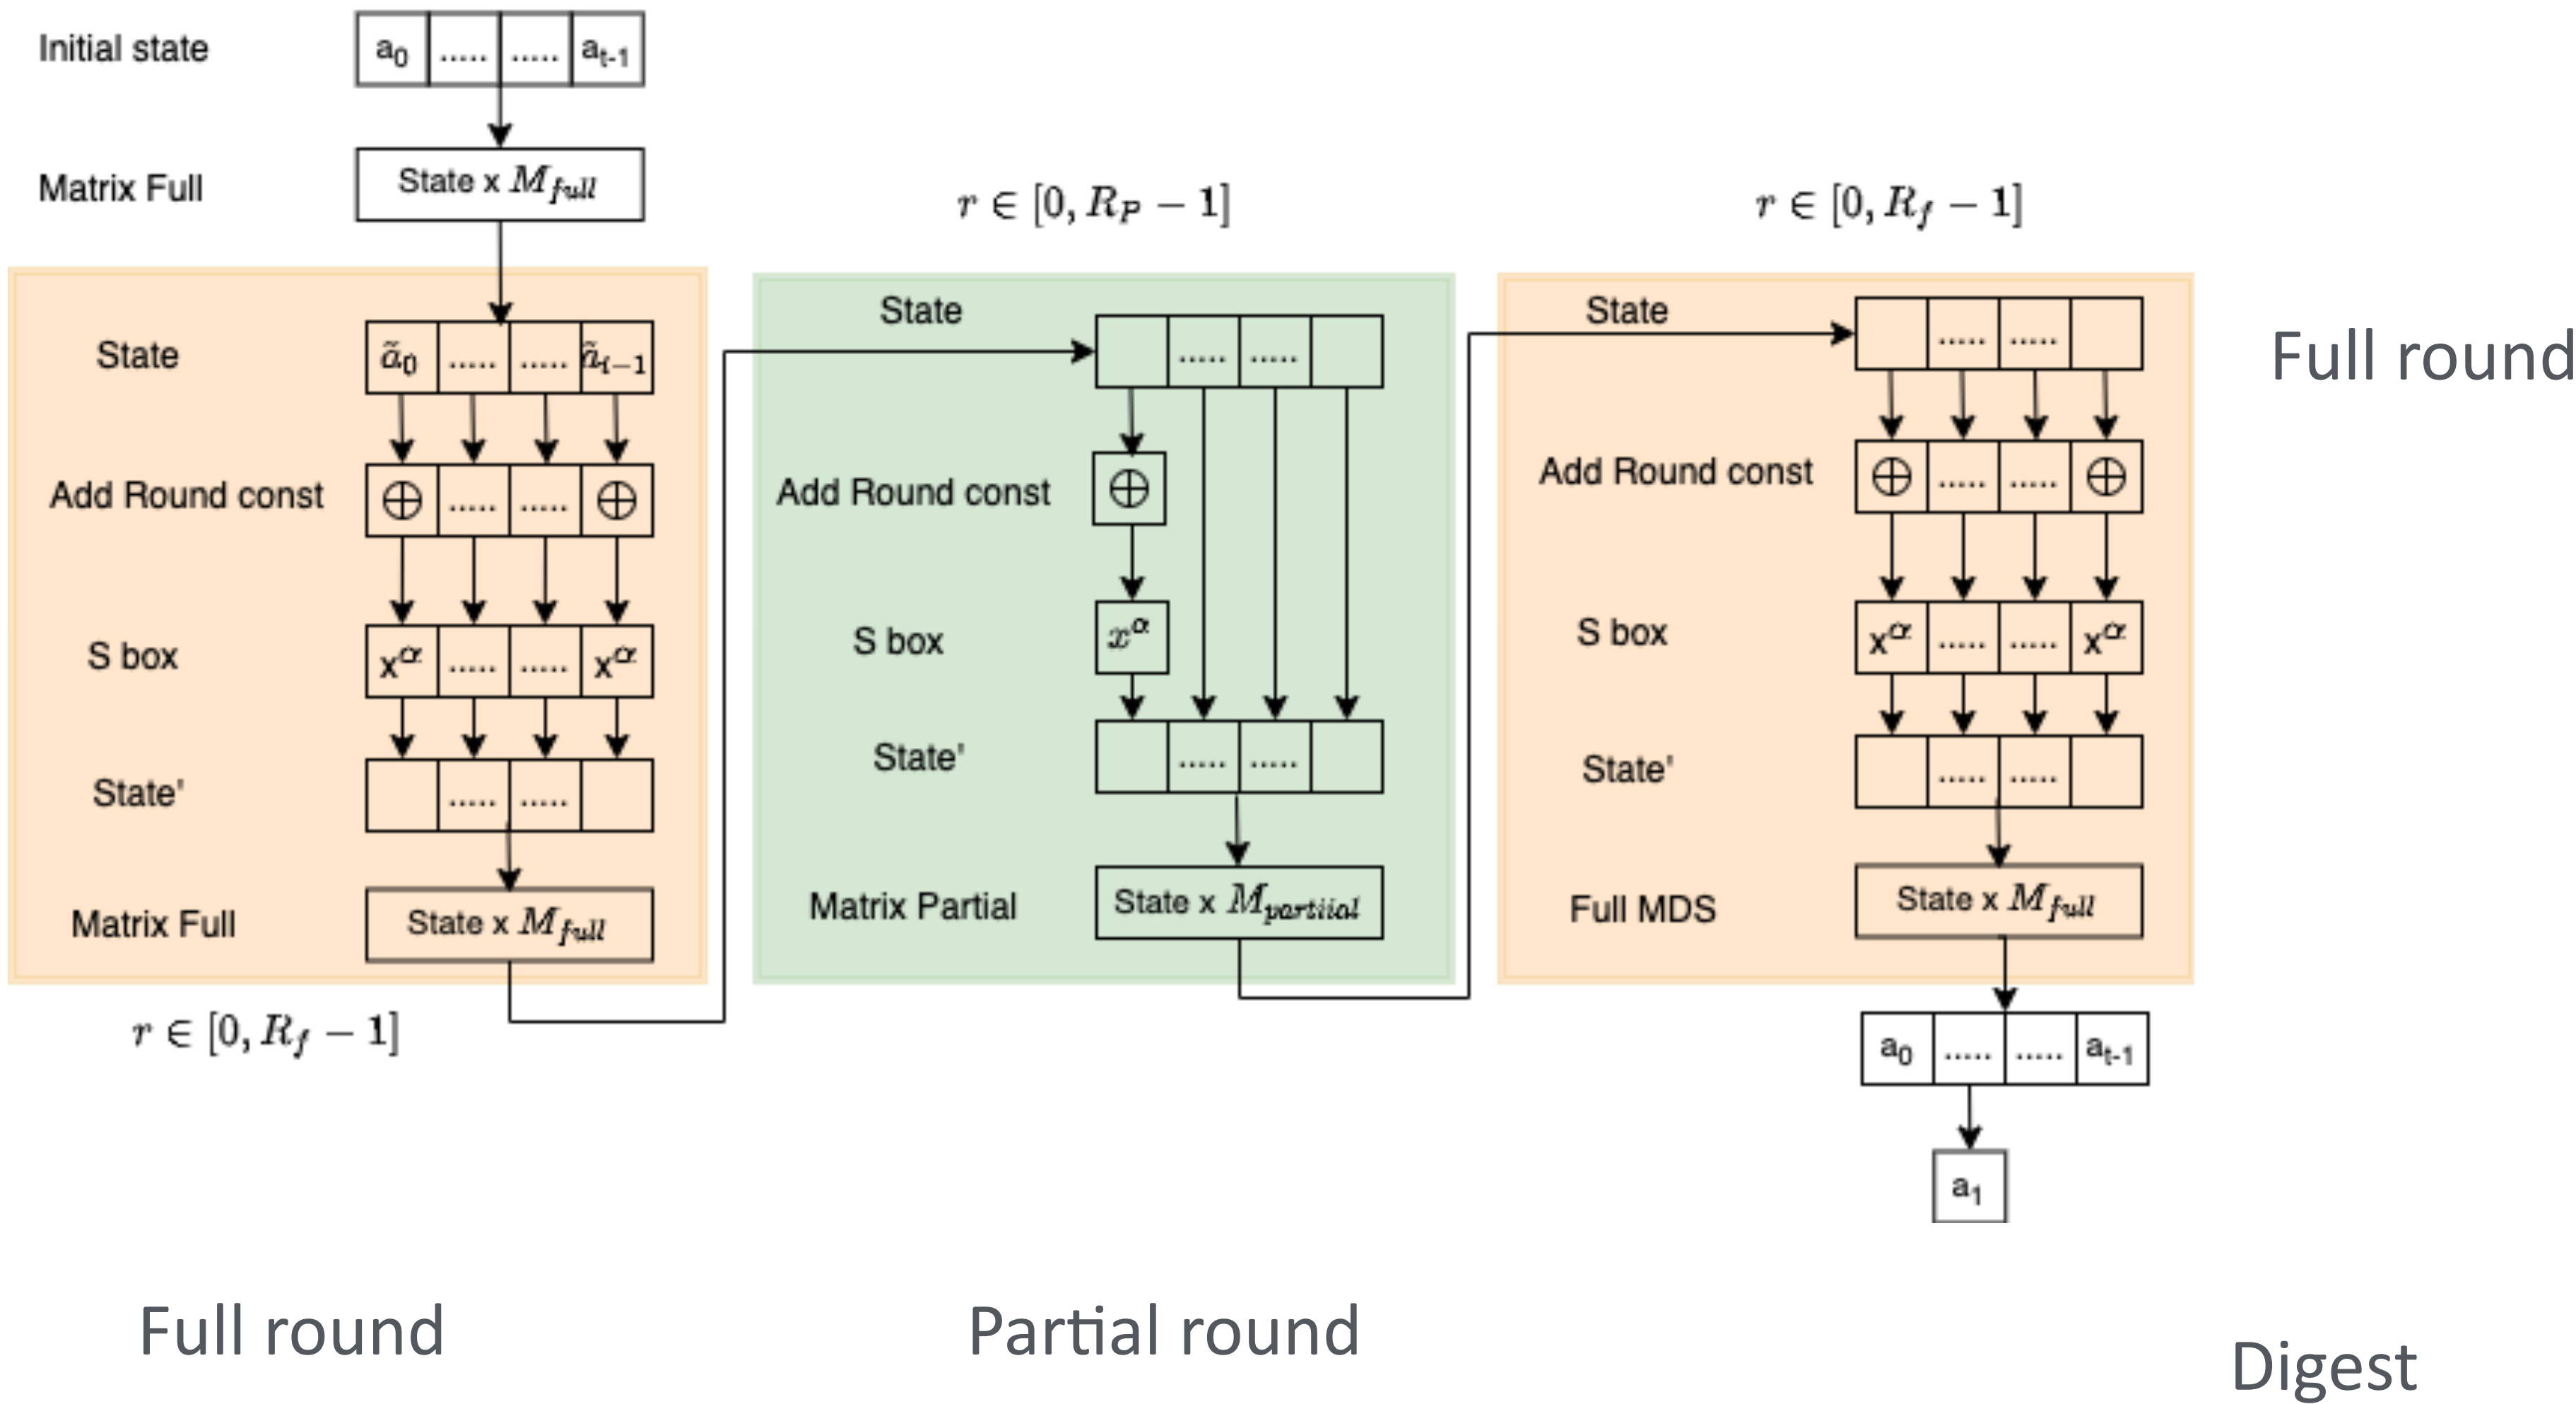
$$state[i] \rightarrow state[i]^\alpha \quad \forall i \in [0, t - 1]$$

- Partial Non-Linear layers:

$$state[0] \rightarrow state[0]^\alpha$$

$$state[i] \rightarrow state[i] \quad \forall i \in (0, t - 1]$$

Poseidon 2



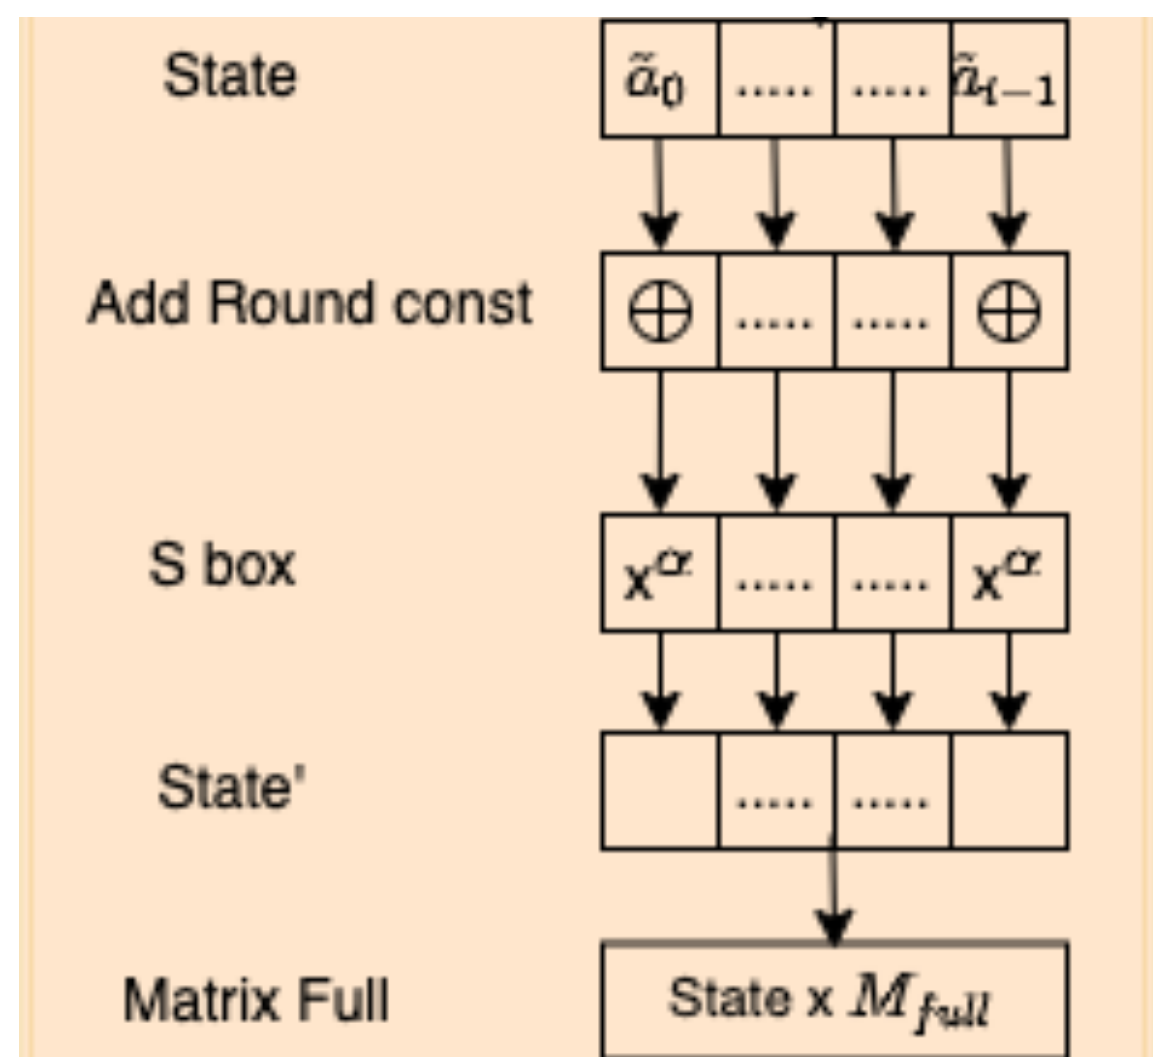
<https://eprint.iacr.org/2023/323>

https://github.com/ingonyama-zk/Poseidon_parameters/tree/master/Poseidon2

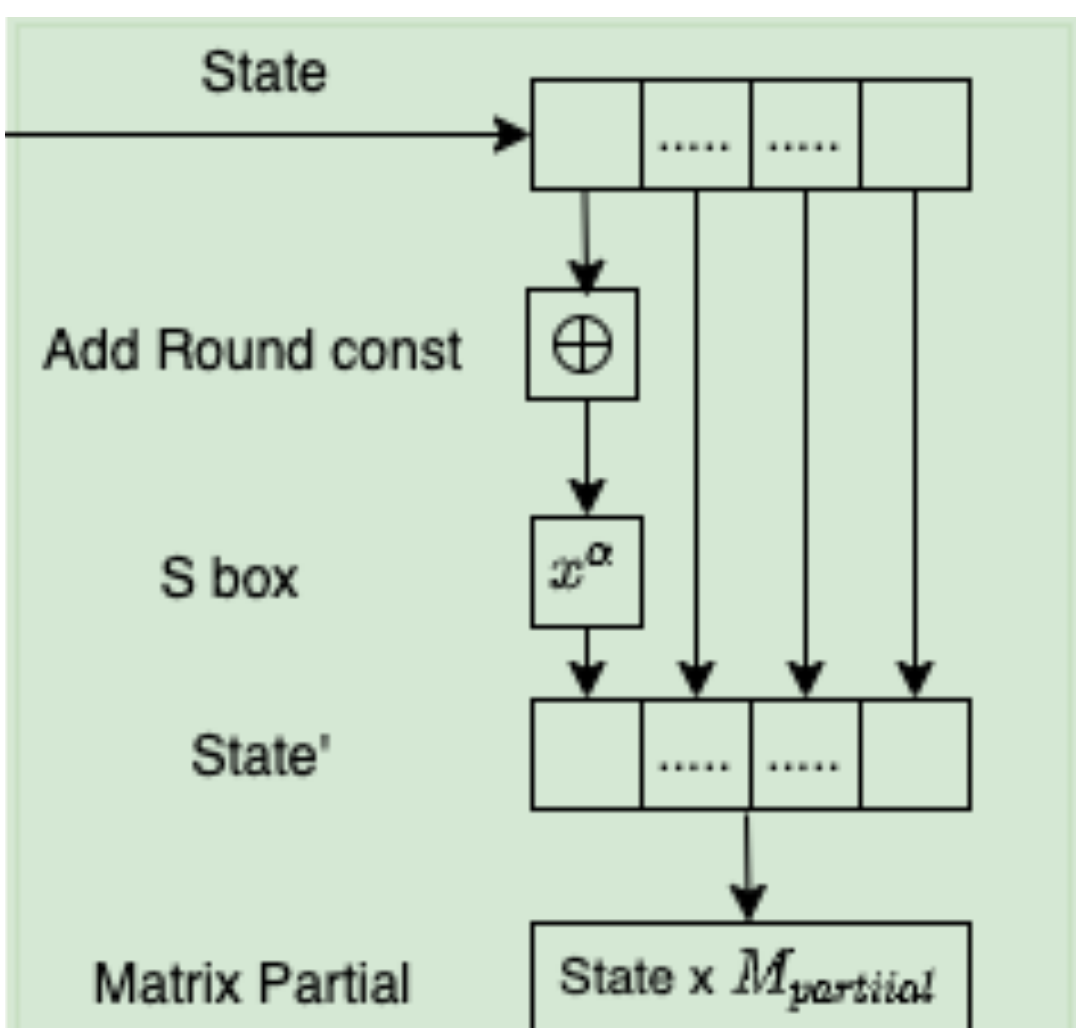
Poseidon 2

- The state is a vector of size $t \in \{2,3,4,\dots,4 \cdot t', \dots 24\}$, with elements in $\mathbb{F}_p^t, p > 2^{30}$

$$M_{full} = \begin{cases} M_4 & \text{if } t = 4 \\ \text{circ}(2 \cdot M_4, M_4, \dots, M_4) \in \mathbb{F}_p^{t \times t} & \text{if } t \geq 8 \end{cases}$$



$$M_4 = \begin{pmatrix} 5 & 7 & 1 & 3 \\ 4 & 6 & 1 & 1 \\ 1 & 3 & 5 & 7 \\ 1 & 1 & 4 & 6 \end{pmatrix}$$



$$\mu_i \in \mathbb{F}_p \setminus \{0,1\}$$

$$M_{partial} = \begin{pmatrix} \mu_0 & 1 & 1 & \dots & 1 \\ 1 & \mu_1 & 1 & \dots & 1 \\ \vdots & \vdots & \ddots & \ddots & \vdots \\ 1 & 1 & \dots & \dots & \mu_{t-1} \end{pmatrix}$$


Poseidon 2: Matrix optimizations

$$M_4 \cdot x = \begin{pmatrix} 5 & 7 & 1 & 3 \\ 4 & 6 & 1 & 1 \\ 1 & 3 & 5 & 7 \\ 1 & 1 & 4 & 6 \end{pmatrix} \cdot \begin{pmatrix} x_0 \\ x_1 \\ x_2 \\ x_3 \end{pmatrix} = \begin{pmatrix} t_6 \\ t_5 \\ t_7 \\ t_4 \end{pmatrix}$$

- Can compute a block in 8 A + 4 M (bit shifts)

$$M_{\text{partial}} = \begin{pmatrix} \mu_0 & 1 & 1 & \dots & 1 \\ 1 & \mu_1 & 1 & \dots & 1 \\ \vdots & \vdots & \ddots & \ddots & \vdots \\ 1 & 1 & \dots & \dots & \mu_{t-1} \end{pmatrix}$$

- μ_i can also be chosen as powers of 2, to turn partial matrix mults to take advantage of bit shifts



Notice: easy arithmetization!

Poseidon 2: Matrix optimizations

$$M_{full} = \begin{cases} M_4 & \text{if } t = 4 \\ \text{circ}(2 \cdot M_4, M_4, \dots, M_4) \in \mathbb{F}_p^{t \times t} & \text{if } t \geq 8 \end{cases}$$

- Circulant matrix multiplications are efficiently performed using Fourier Transform operations.
- But in this case, it is block circulant, and each M_4 multiplication is highly optimized already.
- Using FFT's here is not worth it for the overhead they bring to the computation - not recommended for t sizes commonly used in production.
- However, the MDS matrices themselves are not sacred, and one can find many useful representations for special fields like Goldilocks and M31.